SYSC 5108 W

Deep Learning

Land Use and Land Cover Classification from Satellite Imagery Using Convolutional and Transformer-Based Deep Learning Models

Tan Tran

Student ID: 101058868

Carleton University

April 2026

Abstract

This report presents the development and evaluation of three deep learning models for land use and land cover (LULC) classification from satellite imagery using the EuroSAT benchmark dataset. The work is motivated by satellite-based ghost gear detection efforts at Fisheries and Oceans Canada (DFO), where proprietary operational data cannot be used in an academic setting. EuroSAT, derived from Sentinel-2 imagery provided by the European Space Agency, serves as the public proxy dataset for developing and validating the proposed methods.

Three models are trained and evaluated on a stratified 70/15/15 train/validation/test split of the 27,000-image, 10-class dataset: a Baseline CNN trained from scratch achieving 93.26% test accuracy, a ResNet50 with two-phase ImageNet transfer learning achieving 97.51%, and a Vision Transformer ViT-B/16 with ImageNet-21k transfer learning achieving 98.77%. The ViT-B/16 outperforms both convolutional baselines and exceeds the original EuroSAT benchmark of 98.57% reported by Helber et al. using ResNet50.

Confusion matrix analysis reveals that vegetation-type classes, particularly Permanent Crop, Herbaceous Vegetation, and Pasture present the greatest classification difficulty across all models, accounting for the majority of misclassifications. River confusion with Highway was also identified as a persistent and unexpected error pattern. Results confirm that transfer learning provides a a clear improvement over training from scratch, and that attention-based Transformer architectures outperform convolutional models for satellite scene classification when sufficiently pre-trained.

Table of Acronyms

Acronym	Definition
BN	Batch Normalisation
CNN	Convolutional Neural Network
DFO	Fisheries and Oceans Canada
ESA	European Space Agency
FC	Fully Connected
F1	F1-Score (Harmonic Mean of Precision and Recall)
GPU	Graphics Processing Unit
LULC	Land Use and Land Cover
MHSA	Multi-Head Self-Attention
MLP	Multi-Layer Perceptron
ReLU	Rectified Linear Unit
RGB	Red, Green, Blue
SGD	Stochastic Gradient Descent
ViT	Vision Transformer

Table of Mathematical Notations

Symbol	Description
$x \in \mathbb{R}^{H \times W \times C}$	Input image of height H, width W, C channels
$\hat{y}$	Predicted class label
y	True class label
W_l	Weight matrix of layer l
b_l	Bias vector of layer l
h_l	Output (activation) of layer l
$g(\cdot)$	Activation function
$\mathcal{F}(\cdot)$	Residual mapping function
$\mathcal{L}$	Cross-entropy loss function
η	Learning rate
λ	Weight decay coefficient
$\nabla\mathcal{L}$	Gradient of the loss function
μ	Mean (batch normalisation or channel statistics)
σ	Standard deviation
ϵ	Small constant for numerical stability
γ, β	Learnable scale and shift parameters (batch normalisation)
N	Number of image patches (ViT)
P	Patch size in pixels (ViT)
D	Embedding dimension (ViT)
Q, K, V	Query, Key, Value matrices (self-attention)
d_k	Dimension of key vectors
$\mathbf{E}$	Patch projection matrix (ViT)
$\mathbf{E}_{\text{pos}}$	Positional embedding matrix (ViT)
M	Total number of training examples
K	Total number of classes
p	Dropout probability

Contents

Abstract **1**

Table of Acronyms **2**

Table of Mathematical Notations **3**

1 Introduction **6**
1.1 Motivation . . . 6
1.2 Objectives . . . 7

2 Background **7**
2.1 Related Concepts . . . 7
2.1.1 Deep Neural Networks . . . 7
2.1.2 Convolutional Neural Networks . . . 8
2.1.3 Residual Networks . . . 8
2.1.4 Vision Transformers . . . 8
2.1.5 Transfer Learning . . . 9
2.1.6 Cross-Entropy Loss . . . 9
2.2 Related Work . . . 9
2.3 Rationale for Chosen Methods . . . 10
2.3.1 Why Deep Learning? . . . 10
2.3.2 Why These Three Architectures? . . . 10

3 System Description **10**
3.1 Dataset Description . . . 10
3.2 Neural Network Architectures . . . 11
3.2.1 Baseline CNN . . . 11
3.2.2 ResNet50 with Transfer Learning . . . 11
3.2.3 Vision Transformer ViT-B/16 . . . 12
3.3 Training Process . . . 12
3.3.1 Optimiser and Loss . . . 12
3.3.2 Two-Phase Fine-Tuning for ResNet50 . . . 13
3.3.3 Checkpointing . . . 13
3.4 Testing Process . . . 13

4 Numerical Results **14**
4.1 Performance Metrics . . . 14
4.2 Training and Validation Curves . . . 15
4.3 Confusion Matrices . . . 15
4.4 Top Misclassification Pairs . . . 16
4.5 Model Comparison . . . 16
4.6 Results Discussion . . . 18
4.6.1 Training Curve Analysis . . . 18
4.6.2 Confusion Matrix Analysis . . . 20

4.6.3 Impact of Transfer Learning 21
4.6.4 CNN vs. Transformer 21
4.7 Comparison with Related Results 21

Conclusion **22**

References **23**

1 Introduction

1.1 Motivation

As a public servant employee at Fisheries and Oceans Canada (DFO), we are responsible for supporting the management and protection of Canada's marine and freshwater resources. Among its ongoing mandates is the detection and removal of ghost gear, which are abandoned, lost, or discarded fishing equipment that poses significant ecological risks to marine habitats and species.

The Ghost Gear Problem

Ghost gear is widely recognised as the deadliest form of marine plastic debris. Between 600,000 and 800,000 metric tonnes of ghost gear enter the world's oceans annually [1], and ghost gear accounts for approximately 70% of all macro-plastics in the ocean by weight [2]. Once lost, gear continues to fish indefinitely in a destructive cycle: lost traps and nets retain bait that attracts marine species, which become entrapped and in turn attract further species, repeating until the gear degrades or is retrieved. Animals caught in ghost gear frequently die through suffocation or exhaustion. Ghost gear also damages critical marine habitats including coral reefs and mangroves, and poses hazards to navigation.

The economic consequences are equally significant. Approximately 90% of species caught in ghost gear are of commercial value [2]. A 2021 study found that lobster traps lost in Southwest Nova Scotia alone cost harvesters $175,000 in annual harvest [1]. The cost of traditional mapping efforts to locate ghost gear in three areas of the Canadian Arctic was estimated at $635,000 for a single operation [3]. An AI-based approach, by comparison, has been estimated to cost at most $500,000 annually for year-round regional-scale monitoring, while providing more comprehensive and timely coverage.

Canada's Policy Context

Canada has committed to recovering 100% of all plastics by 2040 and has taken concrete steps toward this goal through the Oceans Protection Plan, Ocean Plastics Charter, Ghost Gear Fund, and Blue Economy Strategy [3]. The Ghost Gear Fund supports harvesters and communities in retrieving lost gear and returning it to the circular economy. Extreme weather events such as Hurricane Fiona have further increased the rate of gear loss, adding urgency to these efforts.

DFO has identified AI-assisted satellite imagery analysis as a key tool for mapping ghost gear at operational scale. In 2023, DFO submitted a Results Fund Proposal in partnership with the Ghost Gear Program to pilot AI solutions for detecting marine litter and predicting gear accumulation densities [3]. Marina Petrovic, Assistant Director of Fisheries Resource Management at DFO, noted that Canada is in a position to build on existing gear-loss datasets by incorporating AI to model potential areas of gear accumulation and to assist in interpreting sonar mapping data, which would allow retrieval efforts to be directed toward areas of highest accumulation or greatest ecological sensitivity [3].

Scope of This Project

DFO currently holds proprietary satellite datasets relevant to this monitoring task; however, this data is subject to confidentiality restrictions that prevent its use in an academic project. To develop and validate deep learning classification methods applicable to satellite imagery, this project uses EuroSAT [4], a publicly available benchmark derived from Sentinel-2 imagery provided by the European Space Agency (ESA). EuroSAT provides a well-established testbed of 27,000 labelled satellite image patches across 10 land use and land cover classes. The underlying task is to classify satellite image patches into distinct categories using deep learning, which applies to the scene-level analysis required for ghost gear detection. Methods validated here can serve as a basis for future application to DFO's restricted operational data.

1.2 Objectives

This project trains and evaluates three deep learning models for Land Use and Land Cover (LULC) classification on EuroSAT:

1. **Baseline CNN**: A compact convolutional network trained from scratch, establishing a lower-bound reference.
2. **ResNet50 with Transfer Learning**: A 50-layer residual network pre-trained on ImageNet [5], fine-tuned using a two-phase strategy.
3. **Vision Transformer ViT-B/16 with Transfer Learning**: A Transformer-based architecture pre-trained on ImageNet-21k [6], fine-tuned on EuroSAT.

The comparative analysis addresses three questions: (i) how much does pre-training improve performance on satellite imagery relative to training from scratch, (ii) do attention-based Transformer architectures outperform convolutional architectures for LULC classification, and (iii) which land use categories are most susceptible to misclassification across all models.

2 Background

2.1 Related Concepts

2.1.1 Deep Neural Networks

A neural network consists of a series of interconnected layers of neurons. Each neuron accumulates a weighted sum of its inputs from the previous layer, adds a bias, and applies a non-linear activation function to produce its output:

$$h_l = g(W_l h_{l-1} + b_l) \tag{1}$$

where W_l are the layer weights, b_l is the bias, and $g(\cdot)$ is the activation function. The most common modern activation function is the Rectified Linear Unit (ReLU):

$$g(z) = \max(0, z) \tag{2}$$

Training adjusts weight values by propagating prediction errors backwards through the network, a process known as backpropagation, to minimise a loss function.

2.1.2 Convolutional Neural Networks

A Convolutional Neural Network (CNN) applies learnable filters across the input using a sliding window, producing feature maps that detect local spatial patterns. Batch normalisation [7] is applied after convolution to stabilise training:

$$\hat{h} = \frac{h - \mu_{\text{batch}}}{\sqrt{\sigma^2_{\text{batch}} + \epsilon}}, \qquad \hat{h}_{\text{scale}} = \gamma \hat{h} + \beta \tag{3}$$

A max pooling operation follows to reduce spatial dimensions and provide translation invariance. Deep CNNs stack many such layers, building increasingly abstract feature representations from low-level edges to complex class-specific patterns.

2.1.3 Residual Networks

Training very deep CNNs is hindered by the vanishing gradient problem. ResNets [5] address this through skip connections:

$$\mathbf{y} = \mathcal{F}(\mathbf{x}, \{W_i\}) + \mathbf{x} \tag{4}$$

where $\mathcal{F}$ is the residual mapping and $\mathbf{x}$ is the identity shortcut. This gives a direct gradient path during backpropagation, allowing stable training of networks with tens to hundreds of layers.

2.1.4 Vision Transformers

The Vision Transformer (ViT) [6] reframes image classification as a sequence modelling problem. An input image is partitioned into $N = HW/P^2$ non-overlapping patches, each linearly projected into a D-dimensional embedding. A learnable class token is prepended and positional embeddings are added:

$$\mathbf{z}_0 = \left[\mathbf{x}_{\text{class}};\ \mathbf{x}_p^1\mathbf{E};\ \ldots;\ \mathbf{x}_p^N\mathbf{E}\right] + \mathbf{E}_{\text{pos}} \tag{5}$$

The sequence passes through Transformer encoder layers applying Multi-Head Self-Attention:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^\top}{\sqrt{d_k}}\right) V \tag{6}$$

The final class token representation is passed to a linear classification head. Unlike CNNs, self-attention operates globally from the first layer, allowing the model to reason about relationships between any two image patches regardless of spatial distance.

2.1.5 Transfer Learning

Transfer learning initialises a model with weights learned on a large source dataset and fine-tunes on the smaller target dataset. The pre-trained weights encode general visual features that transfer across domains, reducing the amount of target-domain data and training time required. Yosinski et al. [10] showed that lower network layers learn highly transferable features while upper layers are more task-specific and benefit most from fine-tuning. This motivates the standard practice of replacing the final classification layer for the target number of classes and fine-tuning the network on the new dataset.

2.1.6 Cross-Entropy Loss

All models are trained using cross-entropy loss:

$$\mathcal{L}(W) = -\frac{1}{M} \sum_{i=1}^{M} \sum_{k=1}^{K} \mathbf{1}[y^{(i)} = k] \log \hat{p}_k^{(i)} \tag{7}$$

where M is the number of training examples, K is the number of classes, and $\hat{p}_k^{(i)}$ is the predicted probability for class k.

2.2 Related Work

The application of transfer learning to fine-tune a deep learning model on DFO ghost gear data has not been previously explored by other data science teams within the department; this would be the first to the author's knowledge. However, we can see that Helber et al. [4] introduced EuroSAT and established a benchmark of 98.57% accuracy using ResNet50, showing that ImageNet pre-trained CNNs transfer well to satellite imagery despite the domain shift from natural images. Marmanis et al. [11] applied ImageNet-pre-trained CNN features to aerial scene classification and showed accuracy gains over models trained from scratch. Cheng et al. [12] confirmed this across multiple remote sensing benchmarks, establishing transfer learning as the standard approach for satellite image classification with limited labelled data.

More recently, Vision Transformers have been applied to remote sensing tasks. Dosovitskiy et al. [6] showed that ViT models pre-trained on large datasets achieve strong performance on image classification benchmarks, with performance generally matching or exceeding CNNs when sufficient pre-training data is available.

2.3 Rationale for Chosen Methods

2.3.1 Why Deep Learning?

Satellite images are high-dimensional data with complex spatial relationships that traditional machine learning approaches with hand-crafted features cannot capture reliably. Deep learning models learn hierarchical feature representations directly from raw pixels, giving better accuracy and generalisation across diverse land cover types.

2.3.2 Why These Three Architectures?

The three models are chosen to form a structured comparison. The Baseline CNN establishes what is achievable from EuroSAT alone. ResNet50 represents the established CNN-based transfer learning approach and provides a direct comparison with the Helber et al. benchmark. ViT-B/16 tests whether global attention-based reasoning provides advantages over local convolutional feature extraction for satellite scene classification, where class identity often depends on spatial layout across the whole image rather than any single local region.

3 System Description

3.1 Dataset Description

The EuroSAT dataset [4] is derived from Sentinel-2 satellite imagery provided by the ESA. It contains 27,000 labelled 64×64 pixel RGB image patches across 10 LULC classes, available at `https://zenodo.org/records/7711810`.

Table 1: EuroSAT Class Distribution

ID	Class	Total Samples	Test Samples
0	Annual Crop	3,000	450
1	Forest	3,000	450
2	Herbaceous Vegetation	3,000	451
3	Highway	2,500	375
4	Industrial	2,500	375
5	Pasture	2,000	300
6	Permanent Crop	2,500	375
7	Residential	3,000	450
8	River	2,500	375
9	Sea / Lake	3,000	450
	Total	**27,000**	**4,051**

A stratified 70/15/15 split is applied, as summarised in Table 2. Stratification ensures each class is proportionally represented in all three subsets.

K-fold cross-validation was considered but not used for this project. First, three architecturally distinct models are trained; with $k = 5$, that requires 15 separate training runs, which

Table 2: Dataset Split

Split	Allocation	Samples
Training	70%	18,900
Validation	15%	4,050
Test	15%	4,051

is computationally prohibitive given the cost of fine-tuning ViT-B/16. In other words, the computational time required does not justify the incremental increase. Second, EuroSAT contains 27,000 images with 2,000–3,000 samples per class, a scale at which a single stratified split produces a test set of over 4,000 samples that is already statistically stable; the reduction in variance from k-fold is minimal at this dataset size.

Per-channel mean and standard deviation are computed from the training split only and applied to normalise all three splits. Augmentations applied to the training split only include random horizontal and vertical flips, random rotation ($\pm 15°$), and colour jitter (brightness and contrast ± 0.2). These transformations are appropriate for top-down satellite imagery which has no canonical orientation. For ViT-B/16, images are resized from 64×64 to 224×224 to match the model's required input resolution. No augmentation is applied to validation or test splits.

3.2 Neural Network Architectures

3.2.1 Baseline CNN

A four-block CNN trained from scratch. Each block follows:

$$\text{Conv2D} \rightarrow \text{BatchNorm} \rightarrow \text{ReLU} \rightarrow \text{MaxPool}$$

with channel widths $\{32, 64, 128, 256\}$ and 3×3 kernels with padding. An adaptive average pooling layer reduces spatial dimensions to 4×4, followed by:

$$\text{Flatten} \rightarrow \text{Linear}(4096 \rightarrow 512) \rightarrow \text{ReLU} \rightarrow \text{Dropout}(p{=}0.4) \rightarrow \text{Linear}(512 \rightarrow 10)$$

The model contains approximately 2.1M parameters and accepts 64×64 RGB input.

3.2.2 ResNet50 with Transfer Learning

ResNet50 [5] is a 50-layer residual network pre-trained on ImageNet. The final fully connected layer, originally configured for 1,000 ImageNet categories, is replaced with a new classification head for 10 EuroSAT classes:

$$\text{Linear}(2048 \rightarrow 256) \rightarrow \text{ReLU} \rightarrow \text{Dropout}(p{=}0.3) \rightarrow \text{Linear}(256 \rightarrow 10)$$

The full model contains approximately 25.6M parameters.

3.2.3 Vision Transformer ViT-B/16

ViT-B/16 [6] divides each 224×224 input into $N = (224/16)^2 = 196$ patches of 16×16 pixels, projected into a $D = 768$-dimensional embedding. The model uses $L = 12$ Transformer encoder layers with 12 attention heads. The pre-trained classification head is replaced with a linear layer mapping the 768-dimensional class token to 10 outputs. The full model contains approximately 86.6M parameters.

Table 3: Model Architecture Summary

Model	**Pre-trained**	**Architecture**	**Parameters**	**Input**
Baseline CNN	No	Convolutional	∼2.1M	64×64
ResNet50	Yes	Conv. + Residual	∼25.6M	64×64
ViT-B/16	Yes	Transformer	∼86.6M	224×224

3.3 Training Process

All models are implemented in PyTorch [13] using the `timm` library [15] for pre-trained checkpoints and trained on a GPU-accelerated environment. The full source code for this project is publicly available on GitHub [14]:

`https://github.com/tranhouse/SYSC5108`

3.3.1 Optimiser and Loss

All models use AdamW [9] with cross-entropy loss. AdamW decouples weight decay from the adaptive learning rate, which gives more effective L2 regularisation than standard Adam. Training parameters are listed in Table 4.

Table 4: Training Hyperparameters

Hyperparameter	**Value**
Loss Function	Cross-Entropy
Optimiser	AdamW
Learning Rate	1×10^{-4}
Weight Decay	1×10^{-4}
Batch Size	32
Epochs	20
LR Scheduler	ReduceLROnPlateau (factor = 0.5, patience = 3)
Gradient Clipping	max norm = 1.0
Random Seed	42

3.3.2 Two-Phase Fine-Tuning for ResNet50

Phase 1 (5 epochs): the backbone is frozen and only the new classification head is trained. This prevents large gradient signals from the randomly initialised head from corrupting pre-trained backbone weights before the head has stabilised. Phase 2 (15 epochs): all layers are unfrozen and fine-tuned end-to-end at 1×10^{-5} (one-tenth the Phase 1 rate), making careful adjustments to backbone features rather than overwriting them.

3.3.3 Checkpointing

At each epoch, if validation accuracy improves, model weights are saved to disk. After training, the best-checkpoint weights are restored for final test evaluation, so the evaluated model is the one that generalised best to the validation set rather than the potentially overfitted final-epoch model.

The training and testing pseudocode are given in Algorithms 1 and 2.

Algorithm 1 General Training Procedure

Require: Training data $(\mathcal{X}, \mathcal{Y})$, learning rate η, epochs N, weight decay λ
Ensure: Trained model parameters θ^*

```
Initialise model parameters θ; best_val_acc ← 0
for epoch = 1 to N do
    for each mini-batch (x, y) in training data do
        ŷ ← model(x; θ)
        L ← CrossEntropy(ŷ, y)
        ∇L ← backprop(L); clip ‖∇L‖ ≤ 1.0
        θ ← AdamW(θ, ∇L, η, λ)
    end for
    for each mini-batch (x, y) in validation data do
        Accumulate validation loss and accuracy (no gradient)
    end for
    Update LR scheduler with validation loss
    if val_acc > best_val_acc then
        best_val_acc ← val_acc; save θ to checkpoint
    end if
end for
Restore θ* from best checkpoint
```

3.4 Testing Process

After training, the best checkpoint is evaluated once on the held-out test set. No augmentation is applied during testing.

Algorithm 2 Testing Procedure

Require: Test data $(\mathcal{X}_{\text{test}}, \mathcal{Y}_{\text{test}})$, trained parameters θ^*
Ensure: Test accuracy, macro F1, confusion matrix

```
Ŷ ← [], Y ← []
for each mini-batch (x, y) in test data do
    ŷ ← arg max model(x; θ*)
    Append ŷ to Ŷ; append y to Y
end for
Compute accuracy, per-class precision/recall/F1, macro F1, normalised confusion matrix
```

4 Numerical Results

4.1 Performance Metrics

Table 5 summarises test set performance across all three models. ViT-B/16 achieves the highest overall accuracy and macro F1, outperforming both convolutional baselines. Transfer learning produces a clear improvement over training from scratch: ResNet50 improves on the Baseline CNN by 4.25 percentage points and ViT-B/16 improves further by an additional 1.26 points.

Table 5: Test Set Performance Summary

Model	**Test Accuracy**	**Macro F1**	**Parameters**
Baseline CNN	93.26%	0.9302	∼2.1M
ResNet50	97.51%	0.9740	∼25.6M
ViT-B/16	98.77%	0.9872	∼86.6M

Table 6 presents macro-averaged precision, recall, and F1 for each model. All three metrics follow the same ranking: ViT-B/16 > ResNet50 > Baseline CNN, and precision and recall are closely matched within each model, indicating no systematic bias toward false positives or false negatives at the macro level.

Table 6: Macro-Averaged Precision, Recall, and F1-Score

Model	**Precision**	**Recall**	**F1**
Baseline CNN	0.9308	0.9311	0.9302
ResNet50	0.9756	0.9728	0.9740
ViT-B/16	0.9877	0.9868	0.9872

Table 7 shows per-class F1 scores for all three models. The weakest classes are consistent across all models: Permanent Crop and River for the Baseline CNN, and Permanent Crop and Pasture for the two transfer learning models. Forest, Industrial, and Residential are the strongest classes across all models, reflecting their visually distinctive signatures at the 64×64 resolution.

Table 7: Per-Class F1-Score Across All Models

Class	Baseline CNN	ResNet50	ViT-B/16
Annual Crop	0.9211	0.9629	0.9833
Forest	0.9802	0.9922	0.9978
Herbaceous Vegetation	0.9001	0.9577	0.9791
Highway	0.9052	0.9813	0.9920
Industrial	0.9685	0.9880	0.9881
Pasture	0.9218	0.9468	0.9780
Permanent Crop	0.8736	0.9441	0.9665
Residential	0.9801	0.9922	0.9955
River	0.8776	0.9800	0.9947
Sea / Lake	0.9738	0.9945	0.9967
Macro F1	**0.9302**	**0.9740**	**0.9872**

4.2 Training and Validation Curves

Figures 1, 2, and 3 show the training and validation loss and accuracy curves for each model.

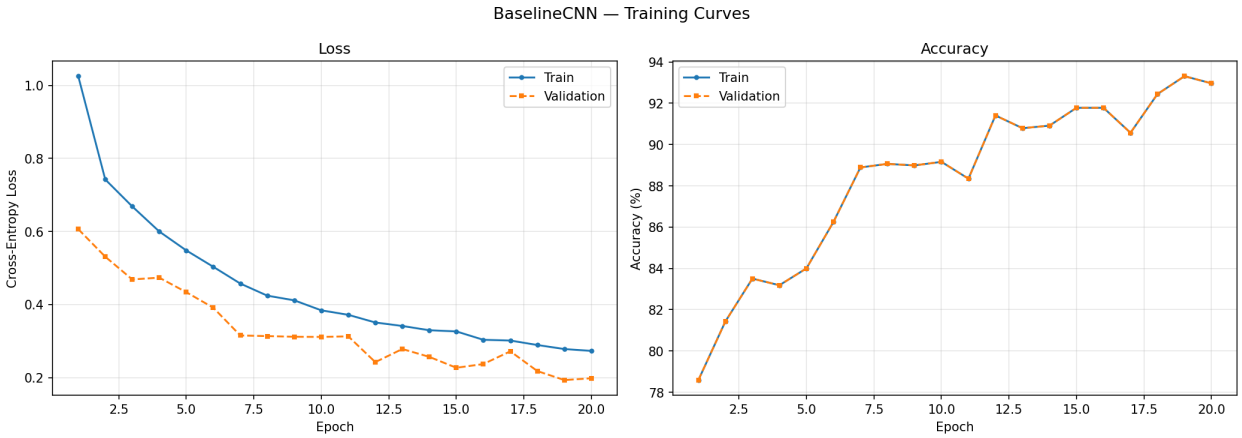

Figure 1: Baseline CNN training and validation loss and accuracy curves.

4.3 Confusion Matrices

Figures 4, 5, and 6 show the row-normalised confusion matrices for each model on the test set. Row normalisation converts counts to per-class recall, making the matrix comparable across classes with different sample sizes. Values on the diagonal represent the proportion of correctly classified samples per class; off-diagonal values indicate confusion between specific class pairs.

Figure 7 presents all three confusion matrices side by side.

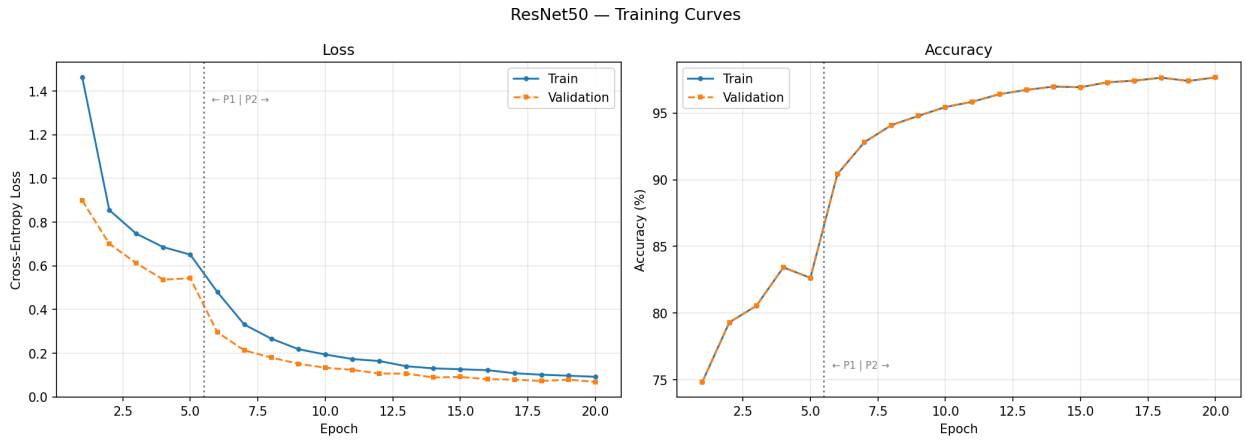

Figure 2: ResNet50 training and validation loss and accuracy curves. The vertical dashed line at epoch 5 marks the Phase 1 / Phase 2 boundary where the backbone is unfrozen for end-to-end fine-tuning.

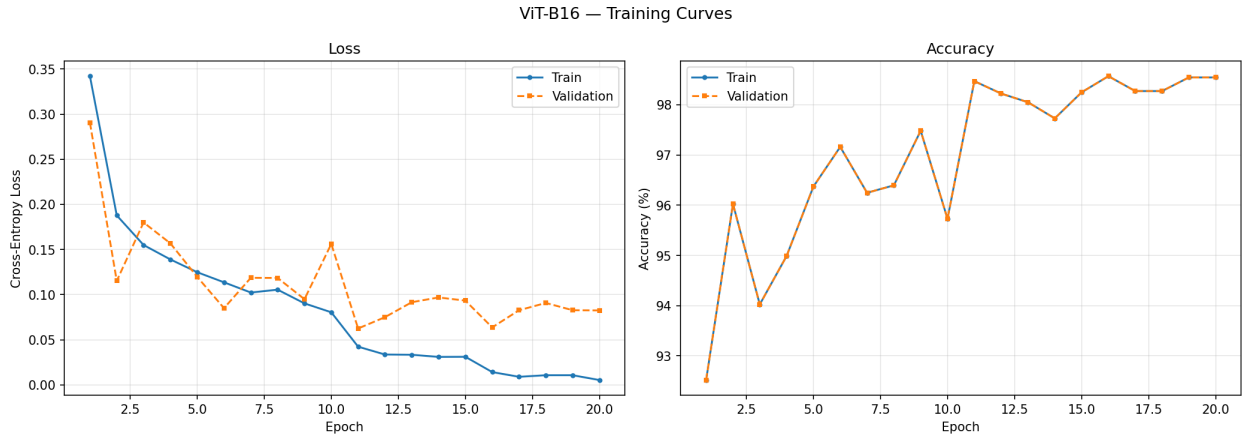

Figure 3: ViT-B/16 training and validation loss and accuracy curves.

4.4 Top Misclassification Pairs

Table 8 lists the ten most common class-pair confusion errors aggregated across all three models. These represent the pairs that were most consistently difficult across all architectures, independent of model-specific behaviour.

4.5 Model Comparison

Figure 8 presents overall test accuracy and macro F1 for all three models in a single bar chart for direct visual comparison.

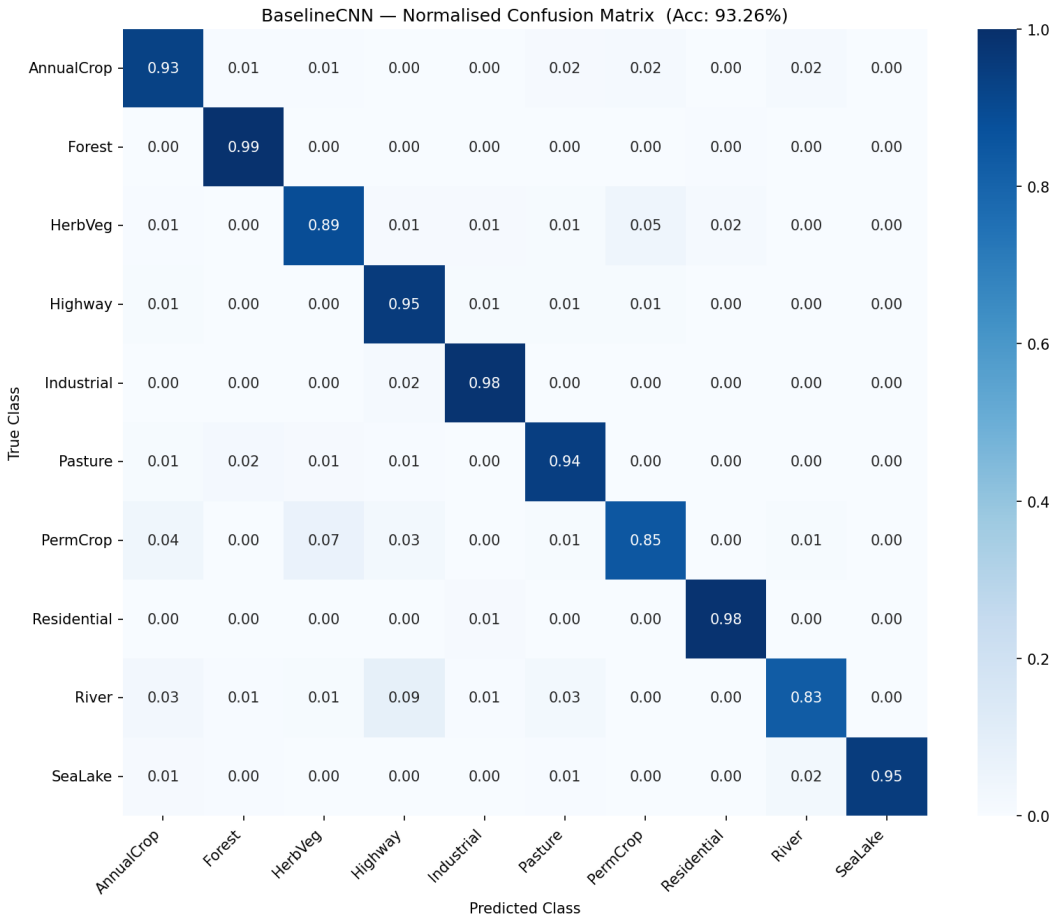

Figure 4: Normalised confusion matrix for the Baseline CNN (93.26%).

Table 8: Most Common Misclassifications (Aggregated Across All Models)

True Class	**Predicted As**	**Count**
Permanent Crop	Herbaceous Vegetation	43
River	Highway	39
Permanent Crop	Annual Crop	32
Herbaceous Vegetation	Permanent Crop	29
Pasture	Herbaceous Vegetation	23
Annual Crop	Permanent Crop	21
Pasture	Annual Crop	16
Permanent Crop	Highway	12
Residential	Industrial	12
River	Annual Crop	12

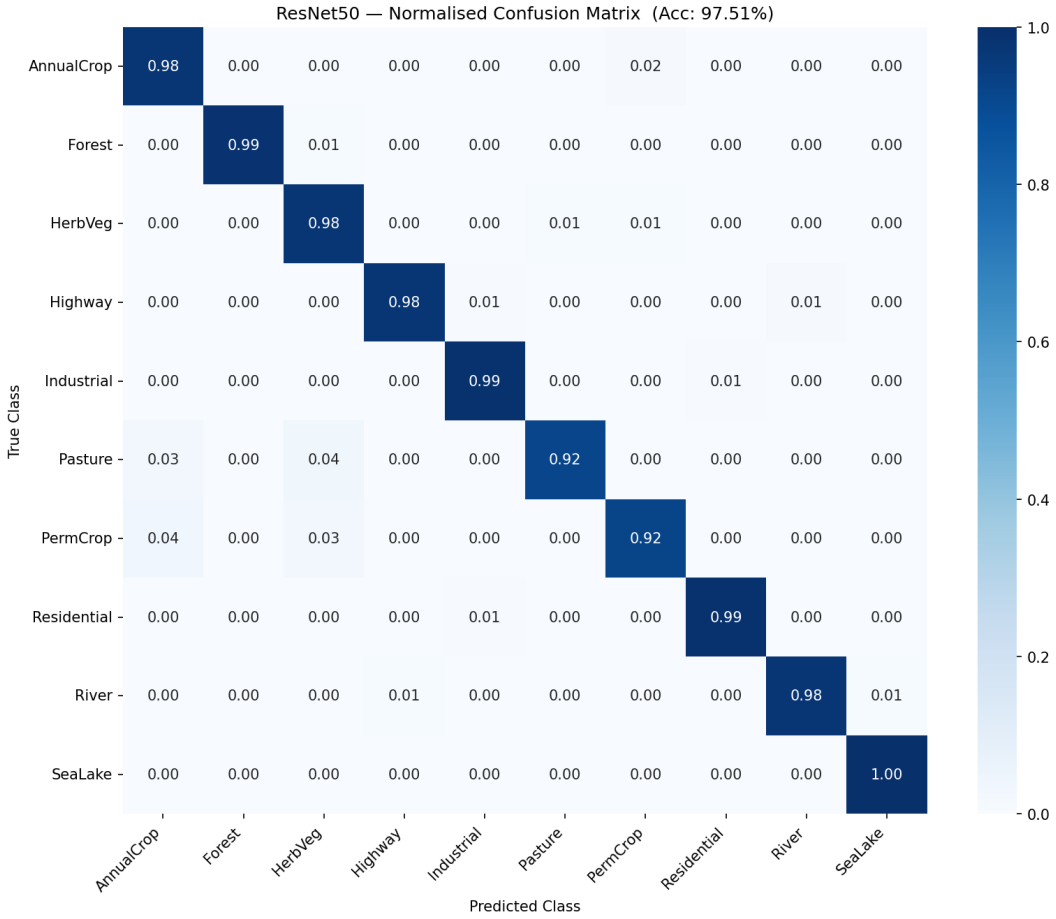

Figure 5: Normalised confusion matrix for ResNet50 (97.51%).

4.6 Results Discussion

4.6.1 Training Curve Analysis

The Baseline CNN curves (Figure 1) show training and validation accuracy tracking closely throughout all 20 epochs, reaching a final validation accuracy of approximately 93%. The small gap between training and validation curves indicates that the model generalises well given its limited capacity, but also that it has reached its performance ceiling. Additional epochs would yield diminishing returns without architectural changes or more data.

The ResNet50 curves (Figure 2) display the characteristic two-phase behaviour of the fine-tuning strategy. During Phase 1 (epochs 1–5, left of the dashed boundary line), validation accuracy rises from approximately 75% to 83% as the classification head is trained against a frozen backbone that immediately provides rich pre-trained feature representations. At epoch 5, when the backbone is unfrozen and the lower learning rate is applied for Phase 2, a sharp improvement occurs and validation accuracy climbs steadily to approximately 97% by epoch 20. Training and validation curves remain closely aligned throughout, indicating effective generalisation.

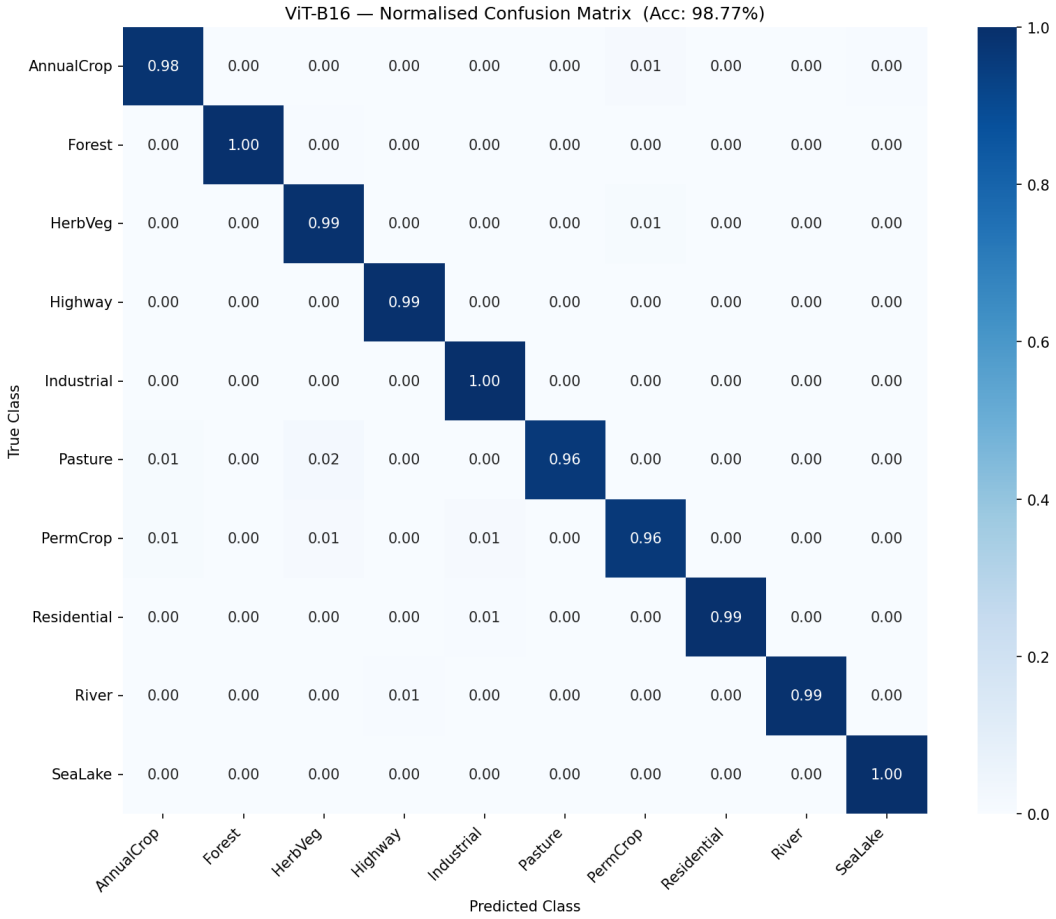

Figure 6: Normalised confusion matrix for ViT-B/16 (98.77%).

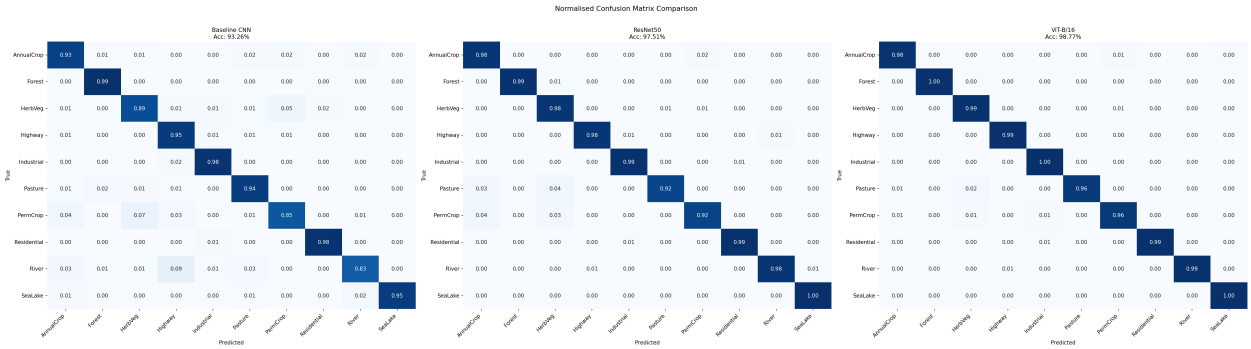

Figure 7: Side-by-side normalised confusion matrix comparison across all three models. The diagonal darkens and off-diagonal values decrease progressively from Baseline CNN to ResNet50 to ViT-B/16, reflecting the accuracy improvements shown in Table 5.

The ViT-B/16 curves (Figure 3) show a different pattern. Training loss descends steeply and approaches near-zero by the final epochs, while validation loss decreases more slowly and with more oscillation. This growing gap between training and validation loss in the later epochs indicates mild overfitting. Nevertheless, the validation accuracy stabilises around

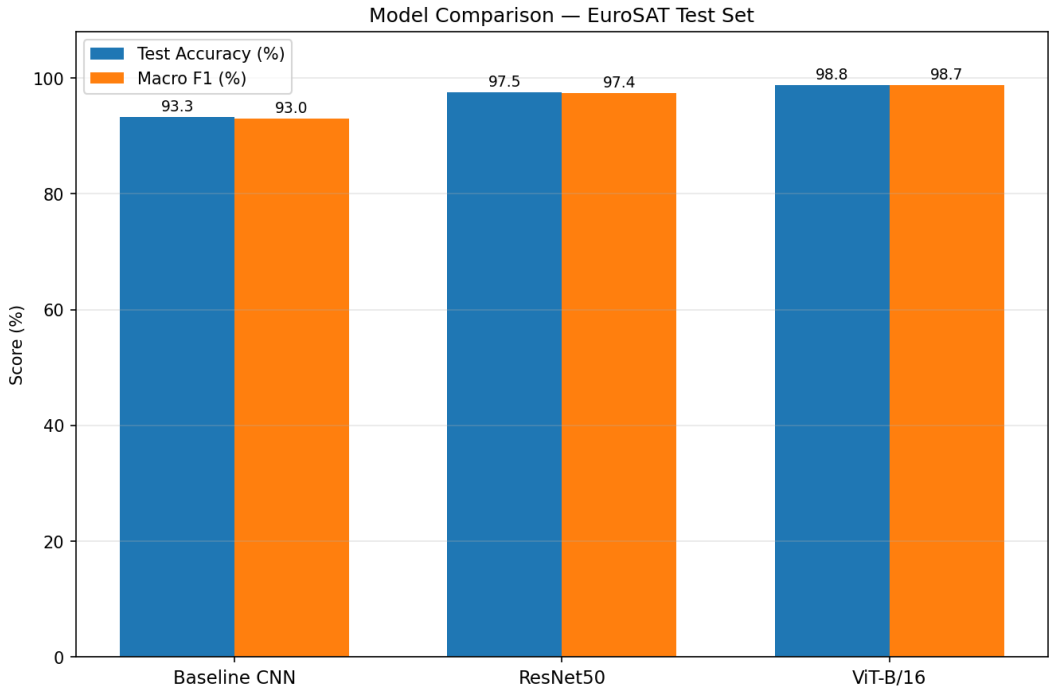

Figure 8: Model comparison - test accuracy and macro F1 on the EuroSAT test set.

98–99%, showing strong generalisation to unseen data despite the training/validation loss gap.

4.6.2 Confusion Matrix Analysis

The confusion matrices reveal several consistent patterns across all models. Forest, Industrial, Sea/Lake, and Residential are the easiest classes to classify, with recall values at or above 0.98 across all three models. These classes have visually distinctive signatures at 64×64 resolution. Forest has dense uniform green canopy, Sea/Lake has open water with characteristic spectral properties, Industrial and Residential have distinct structural textures that even the Baseline CNN can identify reliably.

The vegetation-type classes present the greatest difficulty across all models. Permanent Crop achieves the lowest F1 score in both the Baseline CNN (0.8736) and ResNet50 (0.9441), and is the second-lowest in ViT-B/16 (0.9665). The confusion table confirms this, with Permanent Crop being confused with Herbaceous Vegetation (43 times), Annual Crop (32 times), and vice versa. The assumption is that at the 64×64 resolution and in the RGB colour space, these classes share similar ground cover appearance, such that they are green, textured surfaces without distinctive structural features, and their distinguishing characteristics (crop row spacing, canopy structure) are difficult to resolve at this scale.

The River-to-Highway confusion (39 occurrences) is the most notable unexpected error. Again, at 64×64 resolution, a possibly rationale would be that narrow rivers and roads can appear as similar linear features with comparable widths, especially when water is darker or the image is captured under varying illumination. However, this confusion is largely resolved

by ResNet50 and ViT-B/16, whose pre-trained features appear to encode contextual cues (surrounding vegetation patterns, surface texture) that distinguish water bodies from paved surfaces.

Pasture is consistently confused with Herbaceous Vegetation and Annual Crop, which is expected as all three classes represent open green areas at different stages of agricultural use. The distinction between managed pasture and herbaceous vegetation is subtle even to human observers in small patches.

4.6.3 Impact of Transfer Learning

The 4.25 percentage point accuracy improvement from Baseline CNN (93.26%) to ResNet50 (97.51%) measures the value of ImageNet pre-training on this domain-shifted dataset. The improvement is especially pronounced for classes that were weak in the Baseline CNN; River improves from F1 0.8776 to 0.9800 and Highway from 0.9052 to 0.9813. This suggests that pre-trained features encode texture and structural information that is difficult to learn from 18,900 satellite patches alone.

4.6.4 CNN vs. Transformer

ViT-B/16 outperforms ResNet50 by 1.26 percentage points (98.77% vs. 97.51%), showing that global attention-based reasoning offers an improvement over local convolutional feature extraction for satellite scene classification. The improvement is distributed across nearly all classes, with the most notable gains in Permanent Crop (+2.24 F1 points) and Pasture (+3.12 F1 points), precisely the vegetation classes where spatial layout context is most useful for disambiguation.

4.7 Comparison with Related Results

Table 9 compares results against published benchmarks.

Table 9: Comparison with Related Work

Work	**Year**	**Architecture**	**Dataset**	**Accuracy**
This work	2026	Baseline CNN (scratch)	EuroSAT	93.26%
This work	2026	ResNet50 (transfer)	EuroSAT	97.51%
This work	2026	ViT-B/16 (transfer)	EuroSAT	**98.77%**
[4]	2019	ResNet50	EuroSAT	98.57%
[11]	2016	CNN (transfer)	UC Merced	94.50%
[12]	2017	VGG-VD16 (transfer)	AID	92.07%

The ResNet50 result of 97.51% falls 1.06 points below the 98.57% benchmark reported by Helber et al. [4]. This gap is attributable to differences in training duration, augmentation strategy, and the absence of test-time augmentation in this implementation. Notably, ViT-B/16 at 98.77% exceeds the original EuroSAT ResNet50 benchmark, establishing a new reference result on this dataset using the same public split methodology.

Conclusion

This report trained and compared three deep learning models: a Baseline CNN, ResNet50 with two-phase transfer learning, and ViT-B/16 with transfer learning for LULC classification on the EuroSAT satellite imagery dataset, motivated by ghost gear detection work at DFO.

The experiments produced three clear findings. First, transfer learning provides a substantial benefit on this domain-shifted dataset: ResNet50's 97.51% accuracy represents a 4.25 point improvement over the 93.26% Baseline CNN trained from scratch on EuroSAT alone, showing that ImageNet pre-trained features generalize effectively to satellite imagery. Second, the ViT-B/16 Transformer architecture outperforms the convolutional ResNet50 by 1.26 points, achieving 98.77% and exceeding the original EuroSAT benchmark of 98.57%. This result supports the view that global attention-based reasoning offers advantages for satellite scene classification, where class identity often depends on spatial layout rather than local texture. Third, the most persistent misclassification errors involve vegetation-type classes, particularly Permanent Crop, Herbaceous Vegetation, and Pasture, whose distinguishing characteristics are difficult to resolve at the 64×64 RGB resolution used in this study.

Future work could address these limitations in two directions. Within EuroSAT, the full 13-band multispectral variant includes near-infrared and shortwave-infrared bands with information relevant to distinguishing vegetation types and water bodies, and would likely reduce the vegetation class confusions identified here. More broadly, the pre-trained models and fine-tuning pipeline developed in this project apply to DFO's proprietary satellite data for ghost gear detection, where the trained feature extractors could be further adapted using a small amount of labelled operational imagery.

References

References

[1] Canadian Geographic, "What Lies Beneath: Ghost Gear in Our Oceans," Canadian Geographic Magazine, 2021. [Online]. Available: `https://canadiangeographic.ca/articles/what-lies-beneath-ghost-gear-in-our-oceans/`

[2] Commonwealth Whaling Round Table / International Fund for Animal Welfare (IFAW), "Entangled: Understanding the Full Scope of Ghost Gear," IFAW, 2022. [Online]. Available: `https://www.ifaw.org/international/projects/entangled-whale-rescue-global`

[3] DFO PSDC Team, "Using AI to Detect Ghost Gear," Public Service Data Challenge, Fisheries and Oceans Canada, 2023.

[4] P. Helber, B. Bischke, A. Dengel, and D. Borth, "EuroSAT: A Novel Dataset and Deep Learning Benchmark for Land Use and Land Cover Classification," *IEEE Journal of Selected Topics in Applied Earth Observations and Remote Sensing*, vol. 12, no. 7, pp. 2217–2226, 2019.

[5] K. He, X. Zhang, S. Ren, and J. Sun, "Deep Residual Learning for Image Recognition," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 770–778, 2016.

[6] A. Dosovitskiy, L. Beyer, A. Kolesnikov, D. Weissenborn, X. Zhai, T. Unterthiner, M. Dehghani, M. Minderer, G. Heigold, S. Gelly, J. Uszkoreit, and N. Houlsby, "An Image is Worth 16×16 Words: Transformers for Image Recognition at Scale," in *Proceedings of the International Conference on Learning Representations (ICLR)*, 2021.

[7] S. Ioffe and C. Szegedy, "Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift," in *Proceedings of the 32nd International Conference on Machine Learning (ICML)*, pp. 448–456, 2015.

[8] J. Deng, W. Dong, R. Socher, L.-J. Li, K. Li, and L. Fei-Fei, "ImageNet: A Large-Scale Hierarchical Image Database," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 248–255, 2009.

[9] I. Loshchilov and F. Hutter, "Decoupled Weight Decay Regularization," in *Proceedings of the International Conference on Learning Representations (ICLR)*, 2019.

[10] J. Yosinski, J. Clune, Y. Bengio, and H. Lipson, "How Transferable are Features in Deep Neural Networks?" in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 27, 2014.

[11] D. Marmanis, M. Datcu, T. Esch, and U. Stilla, "Deep Learning Earth Observation Classification Using ImageNet Pretrained Networks," *IEEE Geoscience and Remote Sensing Letters*, vol. 13, no. 1, pp. 105–109, 2016.

[12] G. Cheng, J. Han, and X. Lu, "Remote Sensing Image Scene Classification: Benchmark and State of the Art," *Proceedings of the IEEE*, vol. 105, no. 10, pp. 1865–1883, 2017.

[13] A. Paszke et al., "PyTorch: An Imperative Style, High-Performance Deep Learning Library," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 32, 2019.

[14] T. Tran, "SYSC5108 — Land Use and Land Cover Classification from Satellite Imagery," GitHub repository, 2026. [Online]. Available: `https://github.com/tranhouse/SYSC5108`

[15] R. Wightman, "PyTorch Image Models," GitHub repository, 2019. [Online]. Available: `https://github.com/rwightman/pytorch-image-models`